
The Path-Planning Algorithm A*

System and Device Programming - Project Quer 1 - 2021/22

Giuseppe D'Andrea s303378

Mattia De Rosa s303379



Project Objectives


- Implementation of sequential A^*
- Implementation of two parallel versions of A^*
- Comparing the performance of the different versions

A* Path Finding Algorithm

A* is the most famous “path finding algorithm” and can be used to find the shortest path from a source to a destination in a weighted graph.


- Graph traversal and path finding algorithm
- Extension of Dijkstra’s algorithm
- Uses heuristics to improve performance

Worst-case performance $O(|E|)$



Applications

- Videogames
- Network routing protocols (RIP, OSPF, BGP)
- Maps
- Autonomous vehicles



*An improved A-Star based path planning algorithm for autonomous land vehicles - <https://journals.sagepub.com/doi/full/10.1177/1729881420962263>



Parallelization challenges

- Use of shared resources
- Termination condition
- Nodes re-expansion
- Path reconstruction

Sequential A*



A* components

- Graph G
- Source (s) and destination (d) vertices
- Heuristic function H
- openSet of vertices to be expanded
- closedSet of vertices already explored
- costToCome map of the cost to reach the node



A* iteration

Loop until openSet is empty

1. Take the element n with the lowest f from the openSet
 - a. $f(n) = g(n) + h(n)$
 - b. $g(n)$ is the cost to come
 - c. $h(n)$ is cost to go (heuristic function)
2. Add all neighbors of node n to the openSet if they are not in the closedSet
3. Move n to the closedSet



Heuristic Function

An heuristic function can be

- Admissible

If it never overestimates the goal $h(v) \leq b(v, d)$ with $b(v, d)$ best path between v and d .

- Consistent

If it follows the triangle inequality $h(u) \leq w(u, v) + h(v)$

If the heuristic is admissible and consistent then the first time we reach the destination we have found the shortest path

Hash Distributed A^*



HDA* main concepts

- Vertex ownership

Every vertex is assigned to a processor using an hash function

- Multiple openSets

When expanding a vertex the neighbors are assigned to the owner processor



HDA* issues

- It is impossible to know the order in which operations are completed.
- A higher cost path may be processed before the lowest cost one reaches the openSet.
- We may need to expand a node multiple times.
- This makes detecting a termination condition more difficult.



Termination condition

- A processor can't terminate when his own openSet is empty
 - It may receive more work from another processor
- Solution #1: send a path reached message
 - Only guarantees we reach the destination, but the best path may not be the first found
- Solution #2: terminate only when all open sets are empty
 - Explores many more nodes than needed
 - Requires synchronization (ex: barriers)
- Use both to prune nodes in the open set if they have an higher weight than the current best path

HDA* implementations



HDA* Shared Memory

- openSets and costToCome in shared memory
- Locks to protect from concurrent accesses
- Path reconstruction can be done by a single thread

HDA* Message Passing

- openSets and costToCome in local memory of each thread
- Message queues to send/receive work to/from other threads
- All threads must work together to reconstruct the path

HDA* Termination Condition

- Barriers to synchronize the check of the termination condition
- Every thread checks if its openSet is empty
- If every thread finished its work the search terminates

```
// termination condition
if (openSet.empty()) {
    barrier.arrive_and_wait();
    // set finished flag
    finished[threadId] = openSet.empty();
    barrier.arrive_and_wait();

    // check if all threads finished working, otherwise continue
    if (has_finished()) {
        break;
    } else {
        continue;
    }
}
```





HDA* Synchronization

- `std::barrier` is used in both versions to check for the termination condition
- `std::mutex` is used to protect resources in the shared memory version
- `std::counting_semaphore` is used for path reconstruction in the message passing version

Experimental Results


- K-neighbors graphs
 - 20,000 nodes
 - 50,000 nodes
- City maps exported using OpenStreetMap API
 - Turin: 100K nodes
 - Berlin: 350K nodes
 - New York: 4M nodes



Path reconstruction

- Sequential and shared memory equivalent performance
- Message passing slower than the others

Path Reconstruction Time (μ s) (log scale)



Path reconstruction time (ms)	k-neighbors 20000	k-neighbors 50000	Turin	Berlin	New York
A* Sequential	0.002390	0.003493	0.014338	0.069426	0.403660
HDA* Shared	0.003957	0.005057	0.013706	0.067179	0.399261
HDA* Message Passing	0.476632	1.019622	1.674454	8.170810	44.614082

Results by type of Graph

Graph differences:

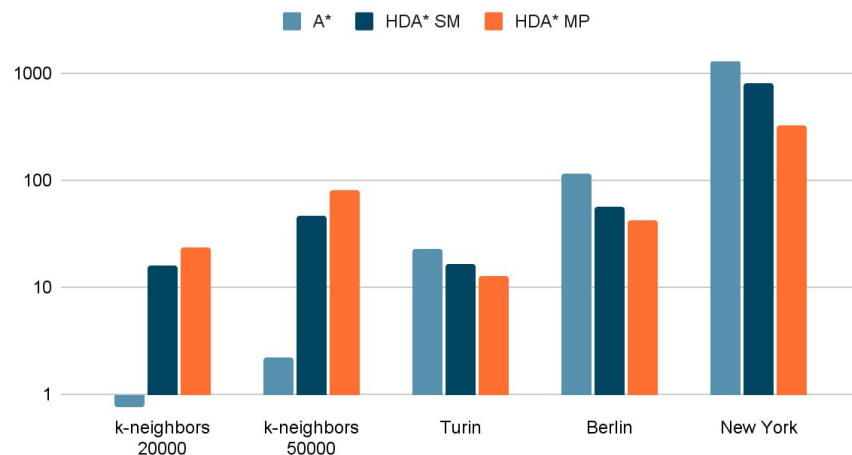
- More edges per node in the k-neighbors graphs
- Weight of edges in the 2 types of graphs

Graph info	k-neighbors 20000	k-neighbors 50000	Turin	Berlin	New York
Nodes	20,000	50,000	95,228	371,857	3,946,582
Edges	78,400	157,0551	105,173	394,062	4,189,184
Edges/Nodes	3.92	31.4	1.1	1.06	1.06


Nodes explored	k-neighbors 20000	k-neighbors 50000	Turin	Berlin	New York
A*	263	409	34,105	137,729	1,028,596
HDA* SM	11,357	18,621	67,363	250,513	3,748,860
HDA* MP	14,499	23,976	59,701	203,157	1,512,746

Results by type of Graph

Execution time (ms) (log scale)

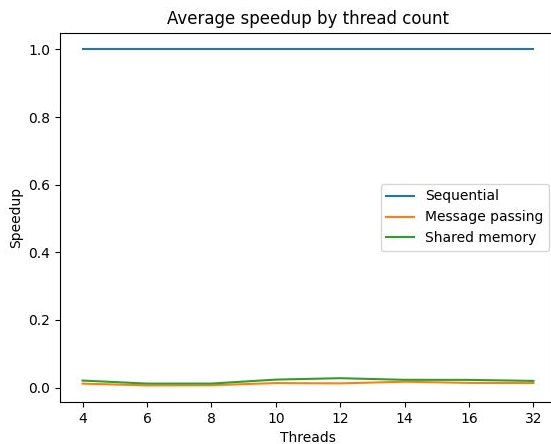


Execution time speedup

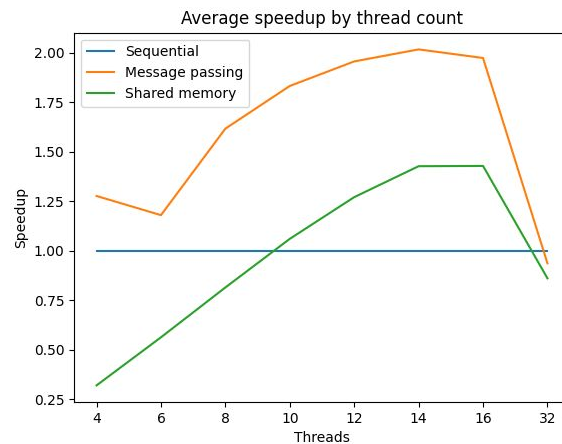


Speedup by thread count

Speedup increases with thread count until we reach the number of logical threads of the processor

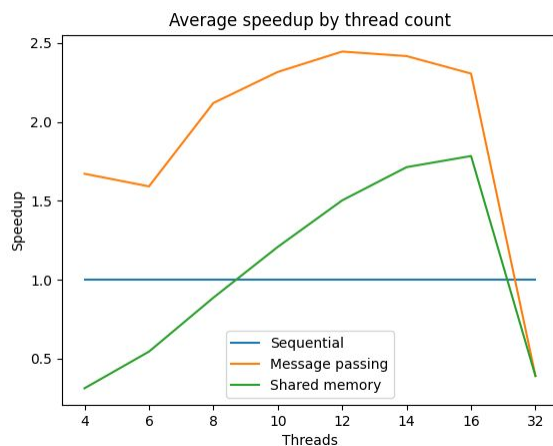


k-neighbors 50000

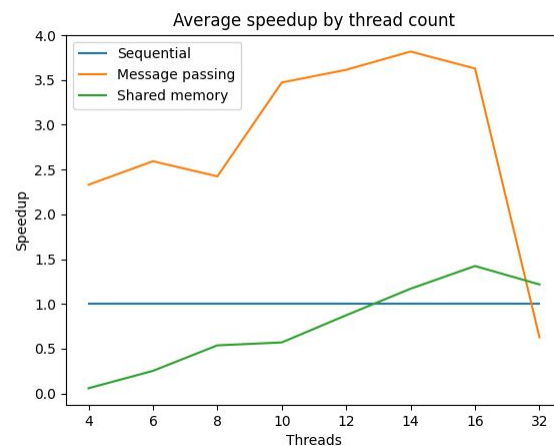


Turin

Speedup by thread count




Berlin



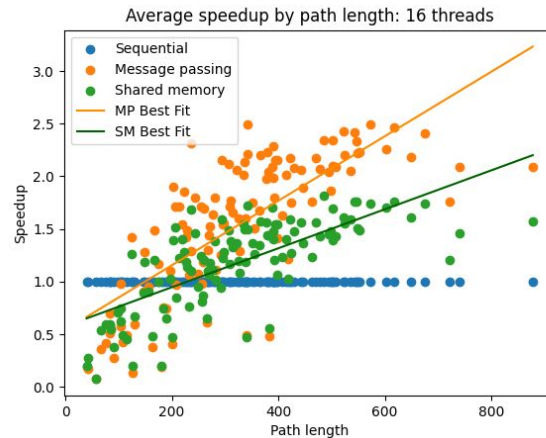
New York

Speedup by path length

Speedup increases on longer and more complex paths





k-neighbors 50000



Turin

Speedup by path length

Speedup doesn't increase for very complex graphs in the shared memory implementation





Conclusions

- Parallel algorithms works better on big graphs where there is a lot of work to distribute
- Message passing performed generally better than the shared memory implementation
- If the heuristic is admissible, consistent and accurate enough to allow to explore a minimal part of the graph, the sequential algorithm could outperform the parallel

Thank you for your attention!
